

CRASH COURSE INTO DOCKER

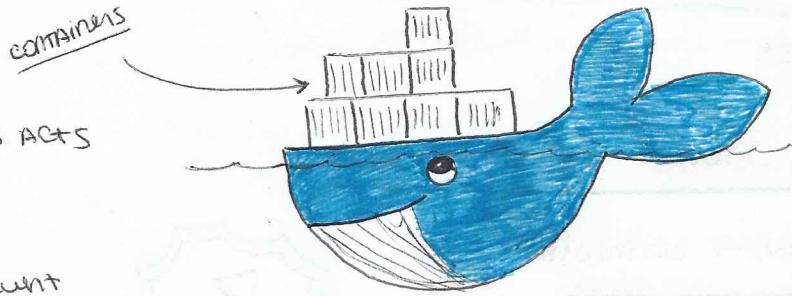
Docker CC

(1)

Docker : TOOL FOR RUNNING APPS IN ISOLATED ENVIRONMENTS

Advantages on VM's

1. Same environment
gives consistency, always acts the same.
2. Sandbox Projects
keeps them secure, prevent any kind of conflicts
3. It works.

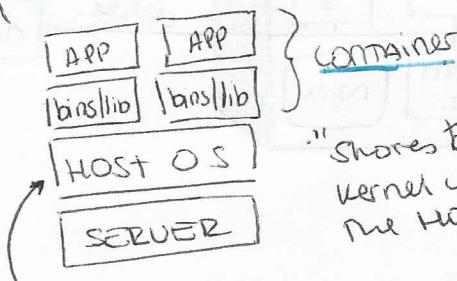


with VM's

VM

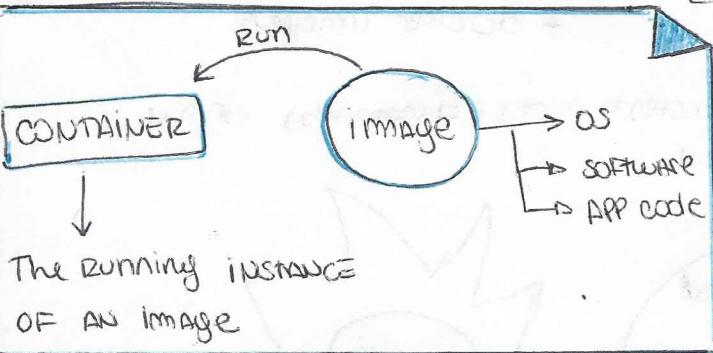


docker run --help



"Shares the kernel with the host"

Ubuntu/debian... All build on the same kernel that Docker uses to build their containers



containers can run specific tasks

docker file



list of steps for creating an image

Image

build



Run

Container

* "for current"



docker run -P port:port, imagename,

-v full path directory,

mount volume.

↳ during dev.

\$ docker build -t imagename, directory.

Dockerfile → <http://hub.docker.com> → get the Dockerfiles

<http://quay.io>

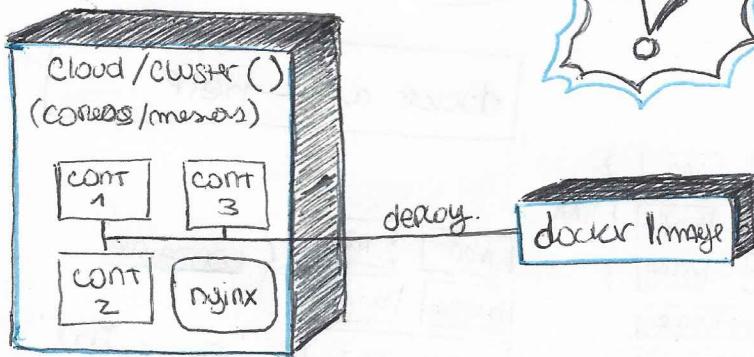
Docker Compose → for running multiple-container APP.

Daemon (computing) is a computer program that runs in background process, rather than being under direct control of an interactive user.

CoreOS/mesos are OS that runs on various computers and shared resources.

How does Docker gets around

- IN cluster computing.



Docker Commands

(BASICS)

- + docker run <image>
Creates a container for that image.
- + docker start/stop, <name|id>
restarts or stops a container
- + docker ps -a
lists all running and stopped (-a) containers
- + docker rm <name|id>
- + docker images

Volumes

Volumes allows docker to ensure persistent parts (directories) of the container even if this last is stopped

Different types

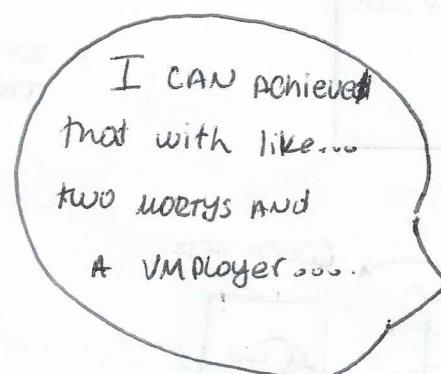
Hosts Volumes

Lives on the Docker host's volume filesystem and it can be accessed from within the container

```
docker run -v /path/on/host:/path/in/container
```

Anonymous Volumes

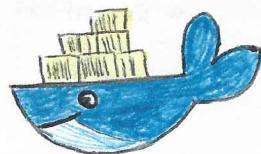
useful when you would rather have Docker handle where the files are stored (it can be difficult to refer to the same volume over the time)



Named Volumes

Ident
Anonymous
but with name

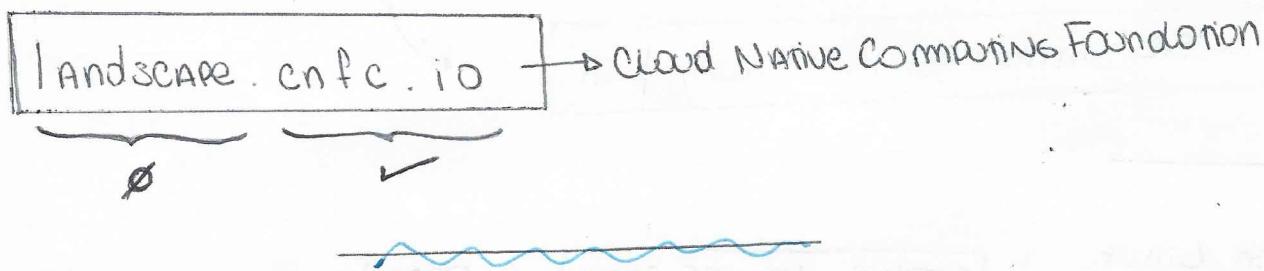
```
docker volume create name
```



(2)

A BIT OF HISTORY OF MIGRATING TECHNOLOGY

- ④ MAINFRAME → PC ← 90's
 - BAREMETAL → VIRTUAL ← 00's
 - DATA CENTER → CLOUD ← 10's
 - HOST → CONTAINER ← Now!
- "Docker migration is a when, not an if"
- "Docker works with sys Admin, devOps and SiteOps"
- #
CONTAINERS ARE CONSISTENT ACROSS THE PLATFORM!
- #
KEEP LESS TIME MAINTAINING AND MORE TIME INNOVATING!



DOCKER INSTALLATION

- DOCKER CE → FREE OPEN SOURCE (Community Edition)
DOCKER EE → ENTERPRISE EDITION

VERSION MATTERS AND ALSO HOW YOU INSTALL IT

3 MAJOR WAYS TO INSTALL DOCKER

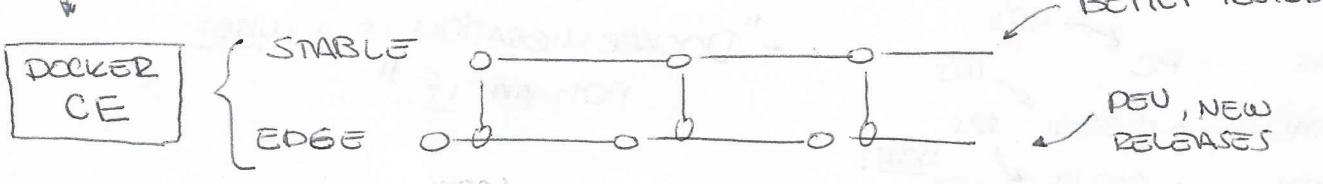
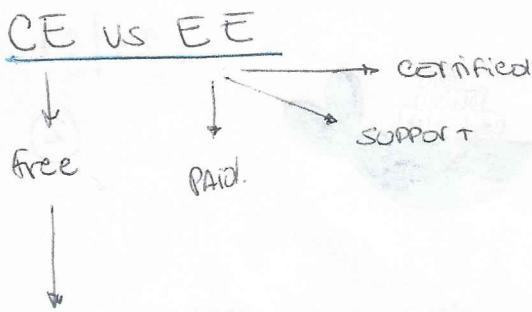
* DIRECT : LINUX (+ per distro), RUNS ON THE KERNEL

* MAC/WIN : SET OF TOOLS THAT LETS YOU RUN DOCKER

MAC OS AND WINDOWS 10 → NATIVELY DOES NOT SUPPORT DOCKER

SO A "VM" MUST BE RUNNING TO ALLOW DOCKER TO RUN ON THESE OS'S

* THE CLOUD : DOCKER FOR GOOGLE CLOUD, AWS, AZURE, ...



INSTALLATION

ON LINUX : # DO NOT USE THE BUILT-IN DEFAULT PACKAGES

DOCKER AUTOMATE SCRIPT TO ADD DOCKER REPOS.

```
curl -sSL https://get.docker.com/ | sh
```

MAGIC!

Docker on Linux

store.docker.com → Details on how to in every distro.

You must use `sudo` for most of docker commands on most distros

CHECK DOCKER VERSION : `sudo docker version`

TWO TOOLS THAT DO NOT COME BUNDLE WITH LINUX

[Docker Machine](#)

[Docker Compose](#)

} BOTH CAN BE INSTALLED, AND MUST BE INSTALLED ON A LINUX DISTRO.
THE HOW-TO ESTÁ EN docs.docker.com.

PERO PARA TENER LA VERSIÓN ULIMA CONVIENE IR A LA PARTE DE 'RELEASES'
DE GITHUB

SECTION 3

Docker Version → Check version and ensure it's working

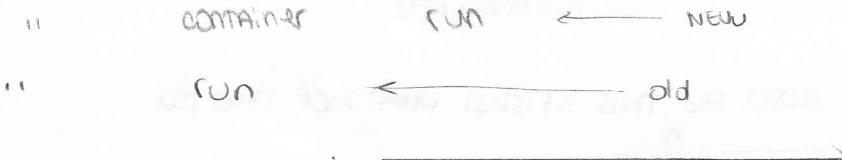
Docker Info → More info and config info.

Docker → --help + commands

Docker Commands Format

New Management Command Format : Docker <command><sub-command>
 +
 <opt>

Docker <cmd> <sub-cmd> <opt>



Docker Container ls

List containers running.

→ ls -a → List all cont.

Image and Containers

Image : is THE APP THAT WE WANT TO RUN

Container : is AN INSTANCE OF AN IMAGE RUNNING AS A PROCESS

You can have multiple cont. of the same image.

→ Starts the cont. with a special cmd in the Docker image

Docker container run -- publish 80:80 nginx

↓ → Starts a new container from that image

The Docker engine looks for an image called "nginx" in Docker Hub, and pulls the latest version

→ The "-- publish 80:80" part of the command expose the port 80 from the local machine and sends all traffic to port 80 on the running container

-- detach → runs container in background and returns the unique container ID

CONTAINER NAMES

IF WE DO NOT SPECIFY A CONTAINER NAME IT WILL GENERATE ONE FROM A LIST OF IT RELATED LIST OF NAMES

`--name [name]` → SPECIFY CONTAINER NAMES

`docker container logs [cont name]` → SHOWS CONTAINER LOGS

`docker container top [cont name]` → PROCESSES RUNNING INSIDE THE CONTAINER

Remove containers

CAN ALSO BE THE 3 FIRST DIGITS OF THE ID.

`docker container rm [c.name1, c.name2, ...]` ↑ ↑
YOU CAN'T REMOVE
RUNNING CONTAINERS

What happens in Docker Container Run

`-f` DO IT'S ANYWAY

#DO NOT CONFUSE CONTAINERS WITH VIRTUAL MACHINES, ARE MORE LIKE PROCESSES RUNNING ON YOUR HOST

CONTAINERS ARE LIKE RESTRICTED PROCESSES RUNNING INSIDE OUR HOST OS

`docker top [cont.name]` → LIST PROCESSES RUNNING INSIDE A SPECIFIC CONTAINER

+ YOU CAN CHECK CONTAINER'S PROCESSES FROM THE HOST. (PS, AUX)

`--env` OR `-e` TO PASS IN VARIABLES, ENVIRONMENT-VARIABLES

`docker container stop` → YOU KNOW
THE DRILL

What's going on with the commands?

`docker container inspect` → DETAILS ABOUT THE CONFIG

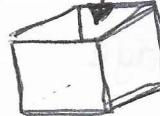
Get a shell inside the container.

Shell

④

docker container run -it

$-i + \underline{-t} = -it$
↳ -interactive



docker container run [options] IMAGE [commands] [args]

This way of running containers lets → after image is specified.

docker container run -it --name proxy IMAGE bash

empty = default command

↳ Prompts a bash shell inside the container.

CONTAINERS RUNS AS LONG AS THEIR FIRST COMMAND RUNS

TO START A CONTAINER AND LAUNCH A SHELL YOU WOULD RATHER USE I-IT THAN RUN, AND START INSTEAD OF RUN BECAUSE YOU CONTAINER ALREADY EXISTS.

docker ~~start~~ container start -it cont.name

→ runs a cmd in a running container.

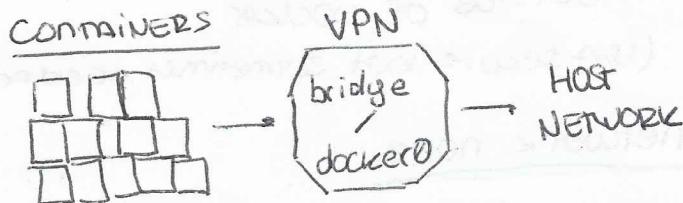
HOW TO JUMP INTO A RUNNING CONTAINER?

docker container exec -it cont.name

Batteries included, but removable

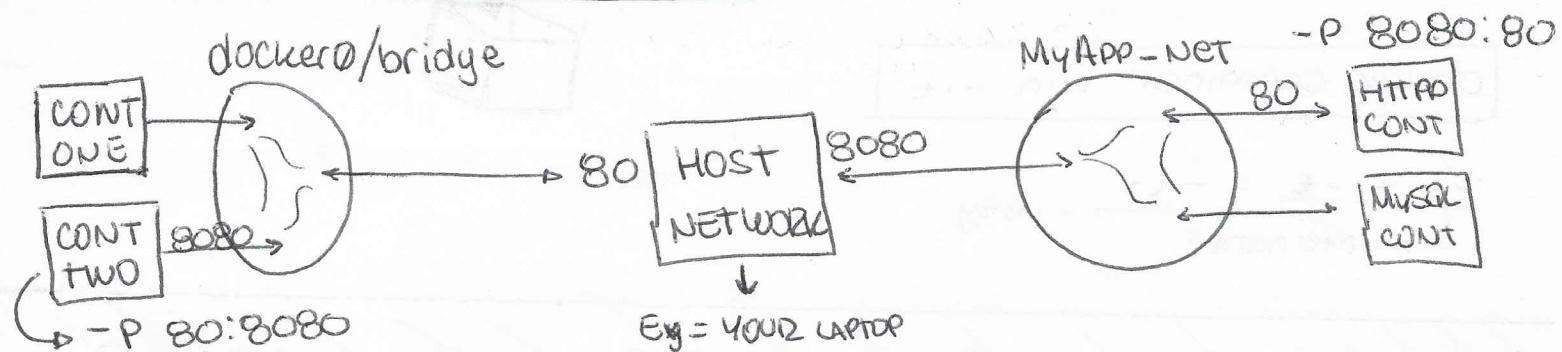
Docker Network : concepts

-p | --publish HOST:CONTAINER
(80:8888)



docker container port cont.name → 80/tcp → 0.0.0.0:80

EJEMPLO DOCKER NETWORKS



CONT ONE AND CONT TWO CAN TALK FREELY OVER THE DEFAULT VPN - docker (it's created by default) BUT CAN'T TALK TO HTTPD CONT OR MYSQL CONT UNLESS IT IS VIA THE HOST.

CLI commands for Networking

`docker network ls` → List Networks.

`docker0` or bridge } connects cont. to local network

`docker network inspect` / + `[networkname]` → Shows details about the network
Also shows info about which containers are attached to that network

--network host

Connects container directly to host interface. Skipping the VIRTUAL INTERFACE OF DOCKER.
(less secure but sometimes needed)

--network none

Removes eth0 and only leaves you with localhost interface in container

Default Networks

- bridge / docker0
- host
- none

CREATE A DOCKER NETWORK

`docker network create my-network-name`

Network Drivers (default: bridge)

- Built in or 3rd party extensions → creates a simple network locally with its own subnet somewhere (172.17.0.0/16)
- that gives you virtual networks features.

TO ATTACH A CONT. TO A SPECIFIC NETWORK:

`docker container run . . . --network net-name`

Connect 3 containers

`docker network connect net, contain.` → ATTACH

`docker network disconnect net, contain.` → DIS-ATTACH

Naming (!) (DNS NAMING)

Docker default sets the hostname to the containers name, but you can also set aliases

Containers should not rely on IP Addresses but on NAMES

It's OK to create custom networks for containers.

`--rm`

`docker container run --rm . . .` → removes / clean up / destroy container when it stops running.

"It is useful for when you need to test something inside a container"

Docker Images

- It is an ordered collection of root file system changes and the corresponding execution parameters for use within a container runtime.
- It is not a complete OS. There is no kernel, kernel modules (e.g. drivers)

Images can ^{be} get from Docker Hub

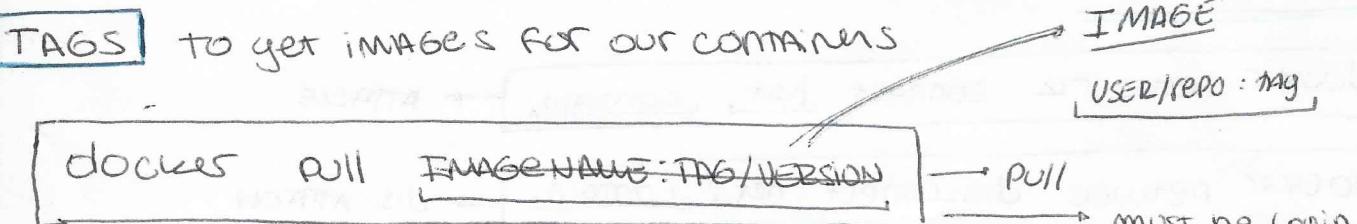
- There is an official signature is ~~an~~ a public

OFFICIAL → "IMAGENAME" → GREAT WAY TO START

NON-OFFICIAL → "IMAGENAME/NAME OF THE REPO"

OFFICIAL IMAGES HAS VERSIONS

TAGS to get images for our containers

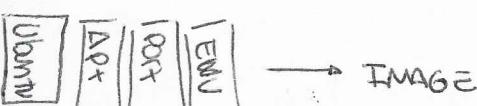


`docker image ls` → SHOWS IMAGES AND + INFO

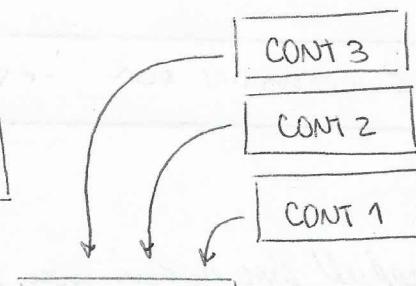
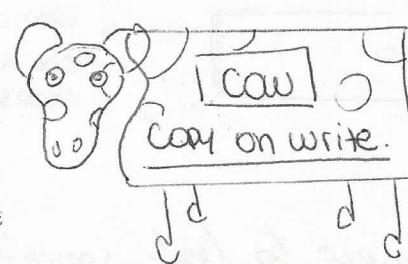
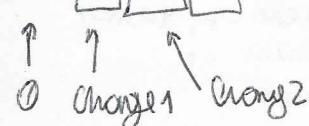
Images History

`docker image history IMAGE NAME` → Shows layer of changes made in an image

Layers (of changes)



NEVER STORE THE SAME IMAGE DATA MORE THAN ONCE IN OUR FILE SYSTEM



CONTAINER IS ONLY THE RUNNING PROCESS

FOR INFO About THE IMAGE

`docker image inspect imagename.`

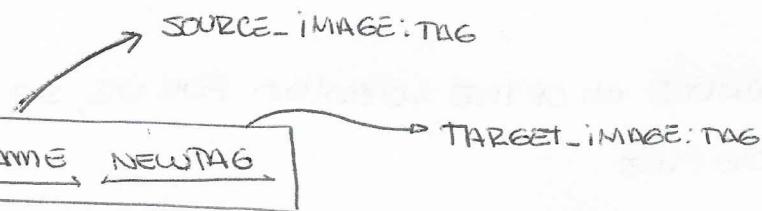
Image Repositories that are official get to own the name of the image without the repo name.

Image TAGGING

`<USER OR ORGANIZATION>/REPOSITORY : TAG` → Identify an Image

it's a pointer to a specific image commit

HOW TO RE-TAG AN IMAGE



HOW TO PUSH

`docker image push`

(must be login to push)

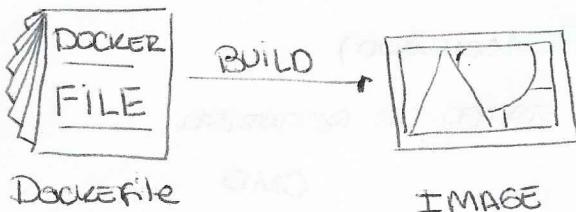
TO LOGIN

`docker login`

TO AUTHENTICATE DOCKER USER

`CAT ~/.docker/config.json`

Building Images



Default name is Dockerfile with caps . "D". (like `MAVENFILE`)

But to build a dockerfile with other name

`docker build -f other-dockerfile`

THE DOCKERFILE

→ USES A LANGUAGE THAT'S UNIQUE TO DOCKER

THE `FROM` COMMAND → MUST BE IN EVERY DOCKER FILE

`FROM debian:release` → PACKAGE manager is the main reason to build cont. FROM UBUNTU, CENTOS, DEBIAN ..

WORDS
TOP ↓ DOWN
ORDER MATTERS

THE RUN COMMAND TI
EXECUTES SHELL COMMANDS!
(UNZIPS, INSTALL PACKAGES, EDIT FILES)

BEST PRACTICE IS TO FEET EACH COMMAND INSIDE ONE LINE WITH "\\" AND &
BECAUSE EACH LINE IS ONE LAYER.

RUN Command-1 \\ → 1 layer
&& comon-2 \\ → 2nd layer
&& comon-3 \\ → 3rd layer
&& comon-4 → 4th layer

LOGS

Docker handle all of the logging for us, so we do not have to log into another log file

All we have to do is redirect to STDOUT and STDERR

RUN ln -sf /dev/stdout /var/stderr/nginx/access.log \\
& ln -sf /dev/stderr /var/log/nginx/error.log

→ HERE WE ARE
REDIRECTING THE
LOGS TO THE ORIGINAL
NGINX LOGFILE

THE EXPOSE COMMAND (OPTIONAL)

By default no port by default is exposed, so we need to specify it here

THE CMD COMMAND (IF NOT ... INHERIT FROM THE FROM COMMAND)

REBUT COMMAND THAT RUNS WHEN A CONTAINER IS STARTED OR RESTARTED.

Building Images

IMAGE NAME
IN THIS DIR ["exec", "cmd", ...]
docker images build -t customnginx !

→ EXECUTE IT BY STEPS, EACH STEP IS A LINE IN THE DOCKERAFILE

Docker CASTED THE STEPS IN THE BUILD USING HASHKEYS SO WHEN YOU MAKE
CHANGES AND RE-BUILD YOU WONT HAVE TO RUN LINES OR STEPS THAT HAVEN'T
BEEN CHANGED

THE WORKDIR COMMAND → RUNS A cd dir.

7

YOU CAN MAKE IT WITH A RUN cd dir BUT THIS IS THE BEST PRACTICE

THE COPY COMMAND

USED TO COPY FILES FROM LOCAL MACHINES/BUILD SERVER TO CONTAINER IMAGES.

Containers Lifetime

SECTION 4

CONTAINERS ARE IMMUTABLE AND Ephemeral

IMMUTABLE INFRASTRUCTURE": JUST RE-DEPLOY, NEVER CHANGE

PERSISTENT DATA → Volumes ①

→ Bind Volumes ②

① Volumes: MAKES SPECIAL LOCATION OUTSIDE OF CONTAINERS UFS

② Bind Volumes: Link container PATH TO host PATH

Volumes And Persistent Data

THE VOLUME command (VOLUME PATH)

ON THE Dockerfile, TELL DOCKER TO CREATE A NEW VOLUME IN THAT PATH.
AND ATTACHED IT IN THE CONTAINER

VOLUMES NEED MANUAL DELETION (JUST FOR ENSURANCE)

TO DO SO:

`docker volume prune`

RUNNING CONTAINERS STORES DATA ON A PATH

("`/var/lib/docker/volumes/` ~ `(-data)`)

AND THEN ROTATED OR MAPPED TO A LOC. IN THE HOST

TO CHECK VOLUMES

TO INSPECT

`docker volume ls`

`docker volume inspect`

Named Volumes

TO SPECIFY A VOLUME NAME. (CREATE)

```
docker container run ... -v NAME:/PATH
```

MORE NAME-FRIENDLY
AND
MORE PATH-FRIENDLY

#

CREATE (NOT IN A RUN COMMAND) (ONLY WAY TO SPECIFY THE DRIVER)

```
docker volume create [OPT] [VOLUME]
```

HOST ALWAYS
WIN

PERSISTENT DATA: Bind Mounting

MAPS A HOST FILE OR DIRECTORY TO A CONTAINER HOST FILE OR DIRECTORY

ON BACKGROUND: IT IS JUST TWO LOCATIONS POINTING TO THE SAME FILE(S)

BIND MOUNTS ARE HOSTS SPECIFIC, SO IT NEED TO BE SPECIFIED ON THE RUN CMD.

How-to

```
... RUN -v /PATH/in/HOST : /PATH/in/cont
```

FULL PATH

CONTAINER
VOLUME

on windows

// DRIVE LETTER

SO, ARE THERE ALWAYS
VOLUMES?

Docker Compose

2 PARTS

① YAML-formatted file which describes our solution for

- Containers
- Network
- Volumes
- ENV VARS
- CONIGS

YAML-code → Config-Language

② A CLI DOCKER-COMPOSE used for local dev/test automation with those YAML files

```
docker-compose --help
```

THE YAML FILE

EASY WAY TO DOCKER CONTAINER RUN

Version :SERVICES # SAME AS DOCKER RUN.

IMAGE

SERVICE NAME : friendlyname (ALSO DNS)

IMAGE :

COMMAND: REPO CMD

ENVIRONMENT:

-e

VOLUMES: -v

SERVICE NAME 2:

VOLUMES: docker volume create

NETWORK: docker network create

ENVIRONMENT:

ONEVAR: var

TWOVAR: var

DOCKER COMPOSE CLI (cool for local dev)

(download it SEPARATELY)

docker-compose up

SETUPS ALL VOLUMES
NETWORKS ETC..
AND STARTS THEM
(it has a -d)

docker-compose down

STOPS ALL CONTAINERS
AND REMOVE CONT/VOL/NET

MORE COMMANDS

top ps ls

docker-compose log

Using Compose to Build Images.

Build Images AT A RUN TIME.

docker-compose build

→ build Images

build VARIABLES (build arguments)

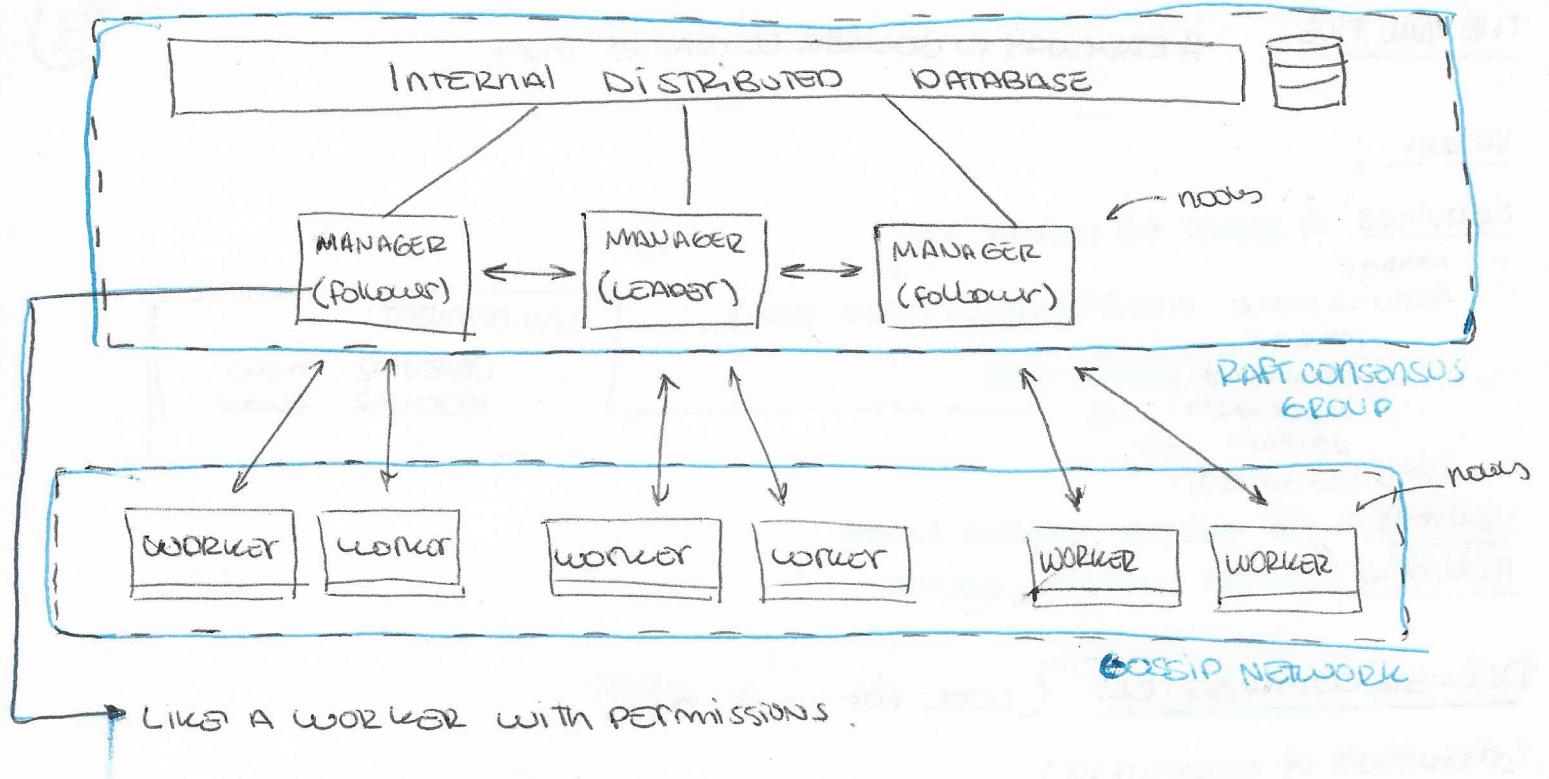
VARS WE DURING BUILDING IMAGES

SWARM INTRO : Built-in orchestrationSECTION 6.SWARM MODE : CLUSTERING SOLUTION INSIDE DOCKER

- WAY OF MANAGING MULTIPLE RUNNING CONTAINERS -

SWARM IS NOT ENABLED BY DEFAULT, ONCE ENABLED:

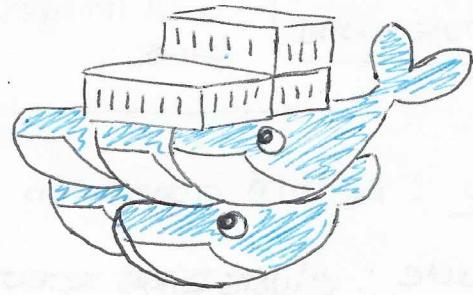
docker { SWARM }



IN A SWARM, we CALL THE DOCKER SERVICE OUR NEW SCALABLE DOCKER RUN



TO KNOW IF SWARM IS ACTIVE



docker info

TO ENABLE SWARM

docker swarm init

docker node inspect

docked node

IS PS

docker node

PROMOTE

DEMOTE

`docker service --help`

CREATE
INSPECT
LOGS
LS
PS
RM
SCALE
UPDATE

The New Docker Run

`docker service ls`

ID	MODE	NAME	REPLICAS	IMAGE
-	-	-	1/1	-

THE GOAL OF THE ORCHESTRATOR IS TO MAKE THESE NUMBERS MATCH

`ps`
`docker service [servname]`

SHOWS THE TASKS OR CONTAINER FOR A SERVICE!
SIMILAR TO `docker container ls` but with nodes

`docker service update <ID> --replicas n`

SCALES A SERVICE SO IT ADDS MORE REPS. (now there are 3 cont. running)

`docker update`

→ limiting and controlling resources without stopping a running container

`docker service update`

→ whole more options,

START AND STOP CONTAINERS

→ The whole point of swarm is to update and make changes without having to take the entire thing down.

SWARM ensures consistent availability

#



MANAGING A 3 NODE CLUSTER

`docker swarm init --advertise-addr`



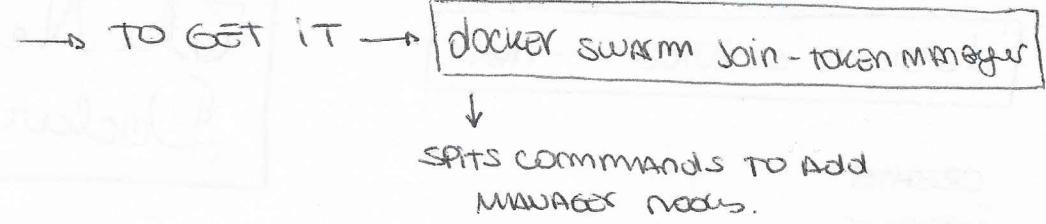
RUN ON NODE 1

SPITS THE

HOST where you have your nodes.

THE JOIN COMMAND

HAS A JOIN-TOKEN



Nodes → HAS CONTAINERS

docker node ls	ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER	STAT
node1	*					

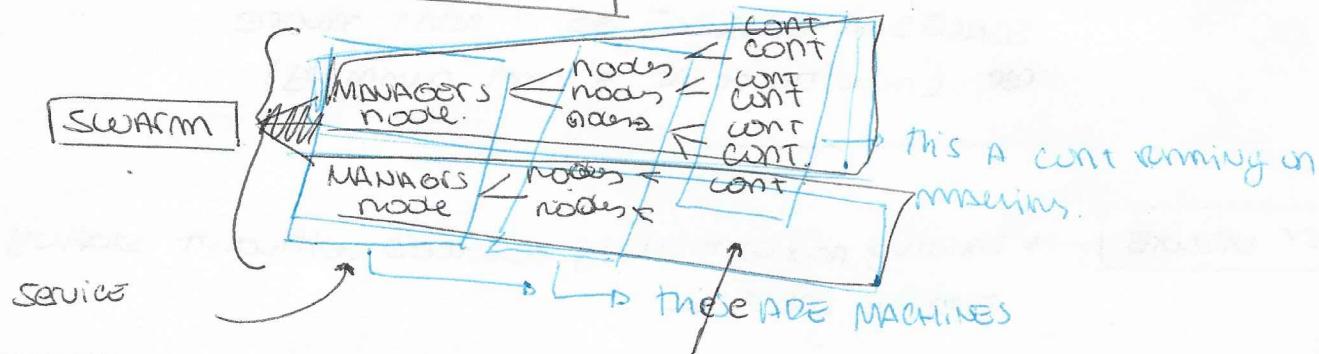
↳ do A container is on node currently talking to this node

NOTE THAT A docker node create doesn't ↳

YOU NEED A MACHINE, each node is a machine

`docker-machine create node_1` → creates a machine

`docker-machine ssh node_1` → connects to that machine

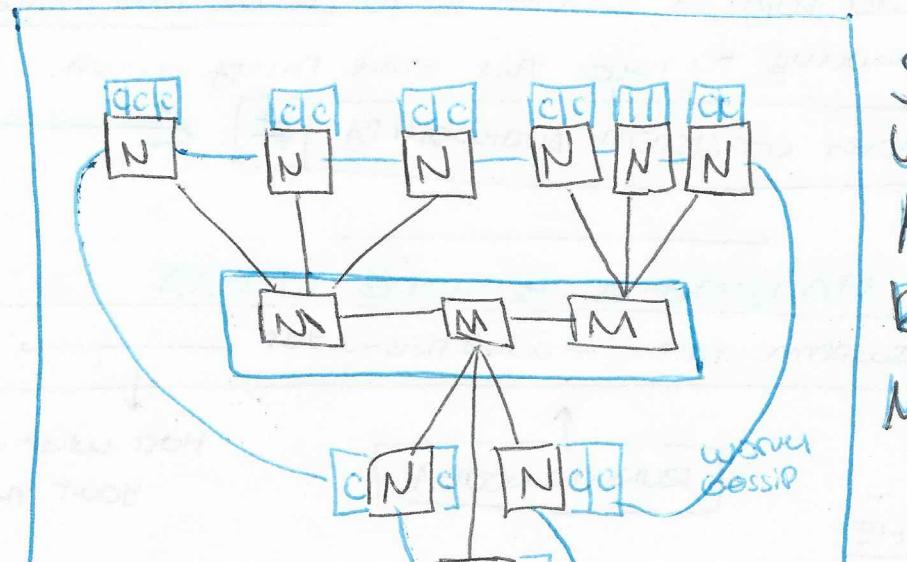


`docker service ps <service>`

THE SERVICE
IS THE NGINX,
.HTTPD, DRUPAL
ETC.

NO MATTER

HOW MANY



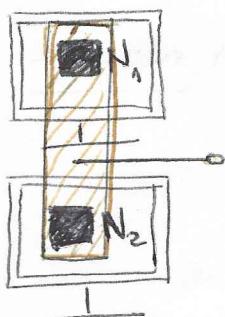
SWARM new NETWORK DRIVER OVERLAY

CREATE A NETWORK AS USUAL
BUT WITH

(10)

-- driver overlay

And it creates a SWARM-wide bridge network.



- ① CONTAINERS ACROSS A NETWORK CAN ACCESS ACROSS THEM, THROUGH what it would be like a VLAN
- ② ONLY FOR INTRA SWARM COMMUNICATION
- ③ OPTIONAL ipsec.
- ④ EACH SERVICE CAN BE CONNECTED TO ≠ NETWORK.

ROUTING MESH

ROUTES incoming PACKETS FOR A SERVICE TO A PROPER TASK

- SPAN ACROSS ALL NODES IN SWARM
- LOAD BALANCES SERVICES ACROSS ALL NODES THEIR TASKS.

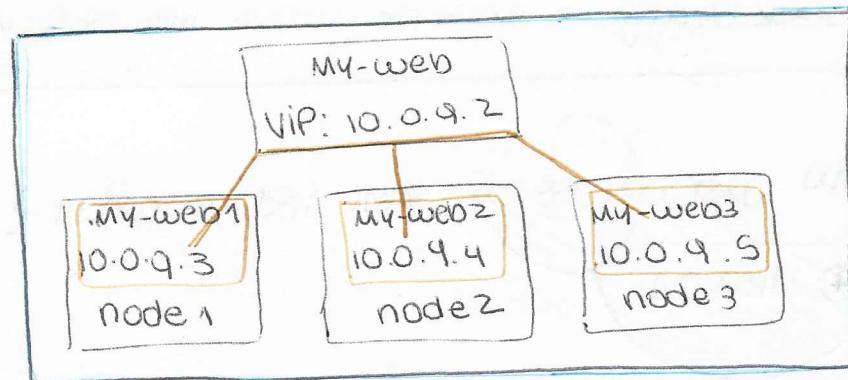
(uses iPVs)

TWO WAYS

CONT - TO - CONT

overlays network
(VIP)

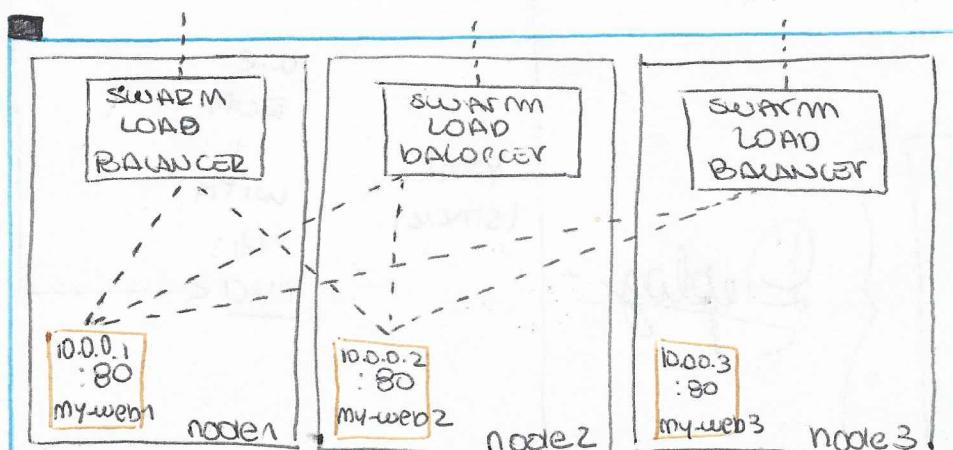
EXTERNAL TRAFFIC INCOMING
(ALL NODES LISTEN)



192.168.99.100

192.168.99.101

192.168.99.102



{ LOAD balancer can do this
because they are on the same network

STACKS: Prod grade compose

Compose files for prod. swarm

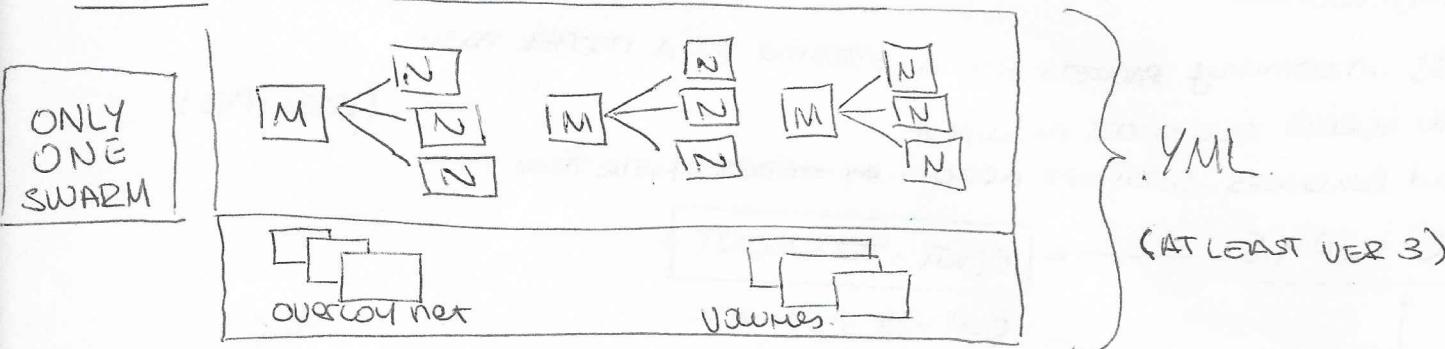
docker stack deploy

Stacks accepts compose files., And do the hard work. Even network.

New deploy: key in Compose file. Can't do build

You don't need docker-compose CLI on Swarm server!

STACK



ON .YML FILE .

docker stack deploy -c example -v0.10.0 -app APP-STACK.YML [Appname]

YOU CAN JUST UPDATE THE .YML AND RE RUN IT !

↓
AT LEAST
-V 3.

.YML + deploy

version: '3'
services:
web:

image: _____
deploy:
replicas:
restart_policy:
condition:
resources:
limits:
cpu:
memory:

.YML
(STACK)

Deploy

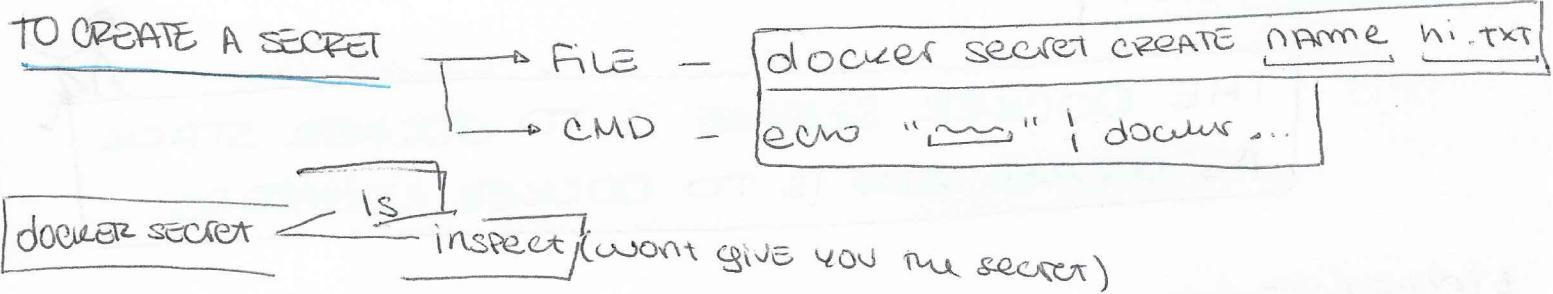
WE
RUN
IT
WITH
THIS
CMD ←

SECRETS IN SWARM

Easiest and secure solution for storing secrets.

SECRET: (SSH keys, TLS cert., sensible DATA)

TO CREATE A SECRET



DOCKER SWARM SECRETS IT'S AN SWARM ONLY THING
BUT DOCKER KNOWS HOW TO MANAGE LOCALLY.

TO MAP A SECRET TO A SERVICE `--secret sec-name`
(SOMETIMES WE USE `-e` TO DO SO TOO)

There ARE ALSO ENVIRONMENT-VAR-FILE SO WE CAN USE A FILE TO A ENV.

`docker service update --secret rm` → REMOVES THE SECRET
BUT RE-DEPLOYS THE SERVICE.

IN ORDER TO USE SECRETS (VERSION: '3.1')

Secrets in Swarm stacks

VERSION: '3.1'

SERVICES:

image: —

secrets:

- PSQL - user
- PSQL - password

environment:

POSTGRES_PASSWORD_FILE: /run/secrets/psql-password
OTHER_ENV_VAR: /run/sec...

② #

ASSIGN SECRETS.
(FIRST ①)

Secrets are in
when stack
is rm

Secrets:

PSQL_USER:

file: ./PSQL-user.txt

PSQL_PASSWORD:

① #

DEFINE ALL
SECRETS AT THE
Bottom

Full App Life cycle → docker-compose for more than one file 

REV

REVELATION: +

THE DOCKER SERVICE IS TO DOCKER STACK
AS DOCKER RUN IS TO DOCKER COMPOSE



A DOCKER SERVICE is ONE OR MORE instances of a SINGLE IMAGE deployed on one or more machines.

A DOCKER STACK is a group of heterogeneous services described in the docker-compose.yml file.

docker swarm is a Clustering and orchestration tool.

UPDATE SWARM

REV

docker stack deploy is THE way of updating a bunch of services.

WHOLE STACK →

docker stack deploy -c file.yml stackname

docker stack def service update --image myapp: newer, servicename

JUST AN IMAGE OF A SERV.

docker service update --env-add NAME=name --Publish-rm 8080

Update PORT (removes) AND Add environment variables ↗

Docker service scale SERV-1=8 SERV-2=6

Updates the amount of replicas of more than one service

Docker Healthcheck

Supported in dockerfile, swarm, compose and run.

STATE

- it exec the command in the container
- this can return only a binary result being exit 0 All Ok! And exit 1 An Error
- Three container state: STARTING, HEALTHY, UNHEALTHY

docker run \

-- health - ^{interval} check = 5s /-- health - end = " " /-- health - retries = 3 /...
.

elastic search : 2

} eg. on docker run.

DOES NOT
REPLACE
MONITORING
TOOL

Image Registries

Default Reg → Docker Hub

Best way of uploading images → Autobuild (within GitHub)

Build on every commit

Docker Cloud

- uses Docker cloud masters and
- bring-your-own-server
- AUTOMATE image build and deploy

^{Hub}
Docker store is for GitHub
what Docker store is for
App Store

3rd Party

QUAY.io → most popular cloud based choice

AWS, Google C., Azure

HAVE THEIR OWN REG.
OPTION.

Docker Registry → PRIVATE GIT SERVER

SELF-HOSTED → GitLab Container Registry